

Synchronization Examples



Operating Systems
Wenbo Shen

Review

- Why we need synchronization?
- Race condition, critical section
- Requirements: ME, Progress, Bounded waiting, Performance
- Locks: acquire, release
 - implementation: test-and-set, compare-and-swap
- Semaphores: wait and signal, implementation

Classical Synchronization Problems

- Bounded-buffer problem
- Readers-writers problem
- Dining-philosophers problem

Bounded-Buffer Problem

- Two processes, the producer and the consumer share **n** buffers
 - the producer generates data, puts it into the buffer
 - the consumer consumes data by removing it from the buffer
- The problem is to make sure:
 - **the producer won't try to add data into the buffer if it is full**
 - **the consumer won't try to remove data from an empty buffer**
 - also call producer-consumer problem
- Solution:
 - n buffers, each can hold one item
 - semaphore **mutex** initialized to the value **1**
 - semaphore **full-slots** initialized to the value **0**
 - semaphore **empty-slots** initialized to the value **N**

Bounded-Buffer Problem

- The producer process:

```
do {  
    //produce an item  
  
    ...  
  
    wait(empty-slots);  
  
    wait(mutex);  
  
    //add the item to the buffer  
  
    ...  
  
    signal(mutex);  
  
    signal(full-slots);  
  
} while (TRUE)
```

Bounded Buffer Problem

- The consumer process:

```
do {  
    wait(full-slots);  
    wait(mutex);  
    //remove an item from buffer  
    ...  
    signal(mutex);  
    signal(empty-slots);  
    //consume the item  
    ...  
} while (TRUE);
```

Readers-Writers Problem

- A data set is shared among a number of concurrent processes
 - readers: only read the data set; they do not perform any updates
 - writers: can both read and write
- The readers-writers problem:
 - allow multiple readers to read at the same time (**shared access**)
 - only one single writer can access the shared data (**exclusive access**)
- Solution:
 - semaphore **mutex** initialized to 1
 - semaphore **write** initialized to 1
 - integer **readcount** initialized to 0

Readers-Writers Problem

- The writer process

```
do {  
    wait(write);  
    //write the shared data  
    ...  
    signal(write);  
} while (TRUE);
```

Readers-Writers Problem

- The structure of a reader process

```
do {  
    wait(mutex);  
    readcount++;  
    if (readcount == 1)  
        wait(write);  
    signal(mutex)  
  
    //reading data  
    ...  
    wait(mutex);  
    readcount--;  
    if (readcount == 0)  
        signal(write);  
    signal(mutex);  
} while(TRUE);
```

Readers-Writers Problem

- The structure of a reader process

```
do {  
    wait(mutex);  
    readcount++;  
    if (readcount == 1) //first reader  
        wait(write); //block write  
    signal(mutex)  
  
    //reading data  
    ...  
    wait(mutex);  
    readcount--;  
    if (readcount == 0)  
        signal(write);  
    signal(mutex);  
} while(TRUE);
```

Readers-Writers Problem Variations

- Two variations of readers-writers problem (different **priority** policy)
 - Reader first
 - no reader kept waiting unless writer is updating data
 - If reader holds data, new reader just moves on and reads
 - writer may starve
 - Writer first
 - once writer is ready, it performs write ASAP
 - If reader holds data, new reader will wait for suspended writer
- Which variation is implemented by the previous code example???

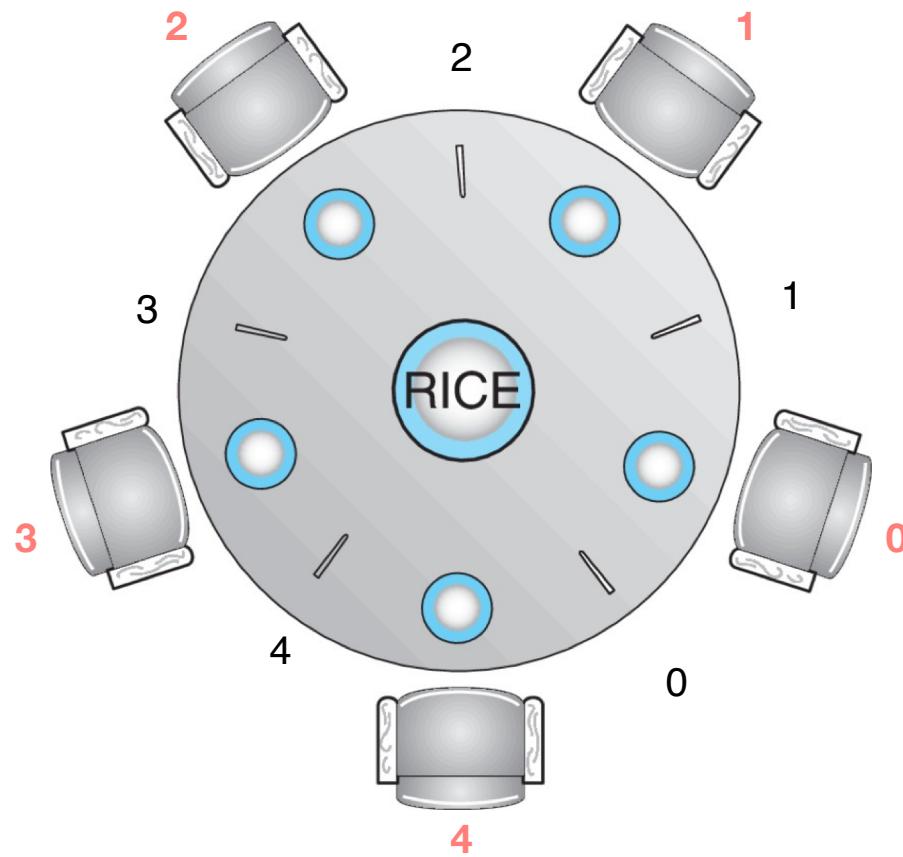
Readers-Writers Problem Variations

- Which variation is implemented by the previous code example???
 - Reader first
- Both variation may have starvation leading to even more variations
 - If writer is in CS and n readers are waiting, one is on **write**, and n-1 are on **mutex**

Dining-Philosophers Problem

- Philosophers spend their lives thinking and eating
 - they sit in a round table, but don't interact with each other
- They occasionally try to pick up 2 chopsticks (one at a time) to eat
 - one chopstick between each adjacent two philosophers
 - need both chopsticks to eat, then release both when done
 - Dining-philosopher problem represents **multi-resource synchronization**
- Solution (assuming **5 philosophers**):
 - semaphore **chopstick[5]** initialized to 1

Dining-Philosophers Problem



Dining-Philosophers Problem

- Philosopher i (out of 5):

```
do  {  
    wait(chopstick[i]);  
    wait(chopStick[(i+1)%5]);  
    eat  
    signal(chopstick[i]);  
    signal(chopstick[(i+1)%5]);  
    think  
} while (TRUE);
```

- What is the problem with this algorithm?
 - **deadlock**

Dining-Philosophers Problem in Practice

```
void *philosopher(void *v)
{
    Phil_struct *ps;
    int st;
    int t;

    ps = (Phil_struct *) v;

    while(1) {

        /* First the philosopher thinks for a random number of seconds */
        ...

        /* Now, the philosopher wakes up and wants to eat.  He calls pickup
           to pick up the chopsticks */
        ...

        pickup(ps);

        ...

        /* When pickup returns, the philosopher can eat for a random number of
           seconds */
        ...

        /* Finally, the philosopher is done eating, and calls putdown to
           put down the chopsticks */
        ...
        putdown(ps);
    }
}
```

Solution 1: do nothing

```
void pickup(Phil_struct *ps)
{
    return;
}

void putdown(Phil_struct *ps)
{
    return;
}
```

```
0 Philosopher 0 thinking for 2 seconds
0 Total blocktime: 0 : 0 0 0 0 0
0 Philosopher 4 thinking for 2 seconds
0 Philosopher 3 thinking for 1 second
0 Philosopher 1 thinking for 2 seconds
0 Philosopher 2 thinking for 1 second
1 Philosopher 3 no longer thinking -- calling pickup()
1 Philosopher 3 eating for 2 seconds
1 Philosopher 2 no longer thinking -- calling pickup()
1 Philosopher 2 eating for 1 second
2 Philosopher 4 no longer thinking -- calling pickup()
2 Philosopher 4 eating for 1 second
2 Philosopher 1 no longer thinking -- calling pickup()
2 Philosopher 1 eating for 2 seconds
2 Philosopher 0 no longer thinking -- calling pickup()
2 Philosopher 0 eating for 1 second
2 Philosopher 2 no longer eating -- calling putdown()
2 Philosopher 2 thinking for 1 second
3 Philosopher 3 no longer eating -- calling putdown()
3 Philosopher 3 thinking for 1 second
3 Philosopher 2 no longer thinking -- calling pickup()
3 Philosopher 2 eating for 2 seconds
3 Philosopher 0 no longer eating -- calling putdown()
3 Philosopher 0 thinking for 2 seconds
3 Philosopher 4 no longer eating -- calling putdown()
3 Philosopher 4 thinking for 2 seconds
```

P2 and p3 cannot
eat at the same time!

Solution 2: A mutex for each chopstick

```
void pickup(Phil_struct *ps)
{
    Sticks *pp;
    int i;
    int phil_count;

    pp = (Sticks *) ps->v;
    phil_count = pp->phil_count;

    pthread_mutex_lock(pp->lock[ps->id]);      /* lock up left stick */
    pthread_mutex_lock(pp->lock[(ps->id+1)%phil_count]); /* lock up right stick */
}
```

```
void putdown(Phil_struct *ps)
{
    Sticks *pp;
    int i;
    int phil_count;

    pp = (Sticks *) ps->v;
    phil_count = pp->phil_count;

    pthread_mutex_unlock(pp->lock[(ps->id+1)%phil_count]); /* unlock right stick */
    pthread_mutex_unlock(pp->lock[ps->id]); /* unlock left stick */
}
```

Solution 2: A mutex for each chopstick

```
0 Total blocktime: 0 : 0 0 0 0 0 0
0 Philosopher 0 thinking for 2 seconds
0 Philosopher 1 thinking for 2 seconds
0 Philosopher 2 thinking for 1 second
0 Philosopher 3 thinking for 2 seconds
0 Philosopher 4 thinking for 1 second
1 Philosopher 2 no longer thinking -- calling pickup()
1 Philosopher 2 eating for 2 seconds
1 Philosopher 4 no longer thinking -- calling pickup()
1 Philosopher 4 eating for 1 second
2 Philosopher 0 no longer thinking -- calling pickup()
2 Philosopher 1 no longer thinking -- calling pickup()
2 Philosopher 3 no longer thinking -- calling pickup()
2 Philosopher 4 no longer eating -- calling putdown()
2 Philosopher 4 thinking for 1 second
3 Philosopher 2 no longer eating -- calling putdown()
3 Philosopher 2 thinking for 2 seconds
3 Philosopher 1 eating for 2 seconds
3 Philosopher 3 eating for 2 seconds
3 Philosopher 4 no longer thinking -- calling pickup()
5 Philosopher 3 no longer eating -- calling putdown()
5 Philosopher 3 thinking for 1 second
5 Philosopher 1 no longer eating -- calling putdown()
5 Philosopher 1 thinking for 1 second
```

Could be deadlock, but ...

Solution 3: Show how deadlock occurs

```
void pickup(Phil_struct *ps)
{
    Sticks *pp;
    int phil_count;

    pp = (Sticks *) ps->v;
    phil_count = pp->phil_count;

    pthread_mutex_lock(pp->lock[ps->id]);           /* lock up left stick */
    sleep(3);
    pthread_mutex_lock(pp->lock[(ps->id+1)%phil_count]); /* lock up right stick */
}
```

```
0 Philosopher 0 thinking for 1 second
0 Philosopher 2 thinking for 3 seconds
0 Philosopher 3 thinking for 1 second
0 Philosopher 4 thinking for 2 seconds
0 Philosopher 1 thinking for 1 second
0 Total blocktime:   0 :   0   0   0   0   0
1 Philosopher 3 no longer thinking -- calling pickup()
1 Philosopher 1 no longer thinking -- calling pickup()
1 Philosopher 0 no longer thinking -- calling pickup()
2 Philosopher 4 no longer thinking -- calling pickup()
3 Philosopher 2 no longer thinking -- calling pickup()
10 Total blocktime:  42 :   9   9   7   9   8
```

Solution 4: An asymmetrical solution

- only odd philosophers start left-hand first, and even philosophers start right-hand first. This does not deadlock.

```
void pickup(Phil_struct *ps)
{
    Sticks *pp;
    int phil_count;

    pp = (Sticks *) ps->v;
    phil_count = pp->phil_count;

    if (ps->id % 2 == 1) {
        pthread_mutex_lock(pp->lock[ps->id]);      /* lock up left stick */
        pthread_mutex_lock(pp->lock[(ps->id+1)%phil_count]); /* lock right stick */
    } else {
        pthread_mutex_lock(pp->lock[(ps->id+1)%phil_count]); /* lock right stick */
        pthread_mutex_lock(pp->lock[ps->id]);      /* lock up left stick */
    }
}

void putdown(Phil_struct *ps)
{
    Sticks *pp;
    int i;
    int phil_count;

    pp = (Sticks *) ps->v;
    phil_count = pp->phil_count;

    if (ps->id % 2 == 1) {
        pthread_mutex_unlock(pp->lock[(ps->id+1)%phil_count]); /* unlock right stick */
        pthread_mutex_unlock(pp->lock[ps->id]); /* unlock left stick */
    } else {
        pthread_mutex_unlock(pp->lock[ps->id]); /* unlock left stick */
        pthread_mutex_unlock(pp->lock[(ps->id+1)%phil_count]); /* unlock right stick */
    }
}
```

Linux Synchronization

- Linux:
 - prior to version 2.6, disables interrupts to implement short critical sections
 - version 2.6 and later, fully preemptive
- Linux provides:
 - **atomic integers**
 - **spinlocks**
 - **semaphores**
 - on single-cpu system, spinlocks replaced by enabling/disabling kernel preemption
 - **reader-writer locks**

Linux Synchronization

- Atomic variables
 - `linux/include/linux/atomic.h`
 - `linux/include/asm-generic/atomic.h`
 - `atomic_t` is the type for atomic integer
- Simple operations
 - `atomic_read()`, `atomic_set()`, ...
- Arithmetic
 - `atomic_{add,sub,inc,dec}()`
 - `atomic_{add,sub,inc,dec}_return{,_relaxed,_acquire,_release}()`
 - `atomic_fetch_{add,sub,inc,dec}{,_relaxed,_acquire,_release}()`

https://github.com/torvalds/linux/blob/master/Documentation/atomic_t.txt

Linux Synchronization

- Bitwise
 - `atomic_{and,or,xor,andnot}()`
 - `atomic_fetch_{and,or,xor,andnot},_relaxed,_acquire,_release{()}`
- Swap
 - `atomic_xchg{,_relaxed,_acquire,_release}()`
 - `atomic_cmpxchg{,_relaxed,_acquire,_release}()`
 - `atomic_try_cmpxchg{,_relaxed,_acquire,_release}()`
- Others
 - Reference count, Misc

https://github.com/torvalds/linux/blob/master/Documentation/atomic_t.txt

Linux Synchronization

- Spinlock
 - [linux/include/linux/spinlock.h](#)
 - Multiple lock/unlock operations supported
- Linux spinlock – x86 implementation

```
ffffffff81a35c00 <_raw_spin_lock>:  
ffffffff81a35c00: 31 c0          xor    %eax,%eax  
ffffffff81a35c02: ba 01 00 00 00  mov    $0x1,%edx  
ffffffff81a35c07: f0 0f b1 17  lock   cmpxchg %edx,(%rdi)  
ffffffff81a35c0b: 75 02          jne    ffffffff81a35c0f <_raw_spin_lock+0xf>  
ffffffff81a35c0d: f3 c3          repz   retq  
ffffffff81a35c0f: 89 c6          mov    %eax,%esi  
ffffffff81a35c11: e9 5a f8 66 ff  jmpq   ffffffff810a5470 <queued_spin_lock_slowpath>  
ffffffff81a35c16: 66 2e 0f 1f 84 00 00  nopw   %cs:0x0(%rax,%rax,1)  
ffffffff81a35c1d: 00 00 00
```

Linux Synchronization

- Linux spinlock – ARM64 implementation

```
fffffff8008a98e78 <_raw_spin_lock>:  
fffffff8008a98e78:    a9bf7bfd      stp    x29, x30, [sp, #-16]!  
fffffff8008a98e7c:    aa0003e3      mov    x3, x0  
fffffff8008a98e80:    d5384102      mrs    x2, sp_el0  
fffffff8008a98e84:    910003fd      mov    x29, sp  
fffffff8008a98e88:    b9401041      ldr    w1, [x2, #16]  
fffffff8008a98e8c:    11000421      add    w1, w1, #0x1  
fffffff8008a98e90:    b9001041      str    w1, [x2, #16]  
fffffff8008a98e94:    d2800001      mov    x1, #0x0          // #0  
fffffff8008a98e98:    d2800022      mov    x2, #0x1          // #1  
fffffff8008a98e9c:    97ff8f53      bl     ffffff8008a7cbe8 <__ll_sc_cmpxchg_case_acq_4>  
fffffff8008a98ea0:    d503201f      nop  
fffffff8008a98ea4:    d503201f      nop  
fffffff8008a98ea8:    35000060      cbnz   w0, ffffff8008a98eb4 <_raw_spin_lock+0x3c>  
fffffff8008a98eac:    a8c17bfd      ldp    x29, x30, [sp], #16  
fffffff8008a98eb0:    d65f03c0      ret  
fffffff8008a98eb4:    aa0003e1      mov    x1, x0  
fffffff8008a98eb8:    aa0303e0      mov    x0, x3  
fffffff8008a98ebc:    97d98c51      bl     ffffff80080fc000 <queued_spin_lock_slowpath>  
fffffff8008a98ec0:    a8c17bfd      ldp    x29, x30, [sp], #16  
fffffff8008a98ec4:    d65f03c0      ret
```

```
fffffff8008a7cbe8 <__ll_sc_cmpxchg_case_acq_4>:  
fffffff8008a7cbe8:    f9800011      prfm   pstl1strm, [x0]  
fffffff8008a7cbec:    885ffc10      ldaxr  w16, [x0]  
fffffff8008a7cbf0:    4a010211      eor    w17, w16, w1  
fffffff8008a7cbf4:    35000071      cbnz   w17, ffffff8008a7cc00 <__ll_sc_cmpxchg_case_acq_4+0x18>  
fffffff8008a7cbf8:    88117c02      stxr   w17, w2, [x0]  
fffffff8008a7cbfc:    35ffff91      cbnz   w17, ffffff8008a7cbec <__ll_sc_cmpxchg_case_acq_4+0x4>  
fffffff8008a7cc00:    aa1003e0      mov    x0, x16  
fffffff8008a7cc04:    d65f03c0      ret
```

Linux Synchronization

- Semaphore
 - [linux/include/linux/semaphore.h](#)

```
struct semaphore {  
    raw_spinlock_t          lock;  
    unsigned int            count;  
    struct list_head        wait_list;  
};
```

- [down\(\)](#), [up\(\)](#)

Linux Synchronization

- Semaphore
 - down()

```
void down(struct semaphore *sem)
{
    unsigned long flags;

    raw_spin_lock_irqsave(&sem->lock, flags);
    if (likely(sem->count > 0))
        sem->count--;
    else
        __down(sem);
    raw_spin_unlock_irqrestore(&sem->lock, flags);
}
EXPORT_SYMBOL(down);
```

- Sleep when holding spinlock?
 - Check __down_common()
 - No, spinlock is released before sleeping

Linux Synchronization

- Semaphore
 - up()

```
void up(struct semaphore *sem)
{
    unsigned long flags;

    raw_spin_lock_irqsave(&sem->lock, flags);
    if (likely(list_empty(&sem->wait_list)))
        sem->count++;
    else
        __up(sem);
    raw_spin_unlock_irqrestore(&sem->lock, flags);
}
EXPORT_SYMBOL(up);
```

Linux Synchronization

- reader-writer locks
 - rw_semaphore
 - RCU (from wiki)
 - **read-copy-update (RCU)** is a synchronization mechanism based on mutual exclusion. It is used when performance of reads is crucial and is an example of space-time tradeoff, enabling fast operations at the cost of more space.
 - Read-copy-update allows multiple threads to efficiently read from shared memory by deferring updates after pre-existing reads to a later time while simultaneously marking the data, ensuring new readers will read the updated data.
 - This makes all readers proceed as if there were no synchronization involved, hence they will be fast, but also making updates more difficult.

POSIX Synchronization

- POSIX API provides
 - mutex locks
 - semaphores
 - condition variable
- Widely used on UNIX, Linux, and macOS

POSIX Mutex Locks

- Creating and initializing the lock

```
#include <pthread.h>

pthread_mutex_t mutex;

/* create and initialize the mutex lock */
pthread_mutex_init(&mutex, NULL);
```

- Acquiring and releasing the lock

```
/* acquire the mutex lock */
pthread_mutex_lock(&mutex);

/* critical section */

/* release the mutex lock */
pthread_mutex_unlock(&mutex);
```

POSIX Semaphores

- POSIX provides two versions – **named** and **unnamed**.
- Named semaphores can be used by unrelated processes, unnamed cannot.

POSIX Named Semaphores

- Creating an initializing the semaphore:

```
#include <semaphore.h>
sem_t *sem;

/* Create the semaphore and initialize it to 1 */
sem = sem_open("SEM", O_CREAT, 0666, 1);
```



- Another process can access the semaphore by referring to its name **SEM**.
- Acquiring and releasing the semaphore:

```
/* acquire the semaphore */
sem_wait(sem);

/* critical section */

/* release the semaphore */
sem_post(sem);
```

POSIX Unnamed Semaphores

- Creating and initializing the semaphore:

```
#include <semaphore.h>
sem_t sem;

/* Create the semaphore and initialize it to 1 */
sem_init(&sem, 0, 1);
```

- Acquiring and releasing the semaphore:

```
/* acquire the semaphore */
sem_wait(&sem);

/* critical section */

/* release the semaphore */
sem_post(&sem);
```

Condition Variables

- `condition x;`
- Two operations are allowed on a condition variable:
 - `x.wait()` – a process that invokes the operation is suspended until `x.signal()`
 - `x.signal()` – resumes one of processes (if any) that invoked `x.wait()`
 - If no `x.wait()` on the variable, then it has no effect on the variable

Condition Variables

- A condition variable is an explicit queue that threads can put themselves on when some state of execution (i.e., some condition) is not as desired (by waiting on the condition);
- some other thread, when it changes said state, can then wake one (or more) of those waiting threads and thus allow them to continue (by signaling on the condition).
- The idea goes back to Dijkstra's use of “private semaphores”;
- a similar idea was later named a “condition variable” by Hoare.

POSIX Condition Variables

- POSIX condition variables are associated with a POSIX mutex lock to provide mutual exclusion: Creating and initializing the condition variable:

```
pthread_mutex_t mutex;  
pthread_cond_t cond_var;  
  
pthread_mutex_init(&mutex,NULL);  
pthread_cond_init(&cond_var,NULL);
```

POSIX Condition Variables

- Thread waiting for the condition $a == b$ to become true:

```
pthread_mutex_lock(&mutex);
while (a != b)
    pthread_cond_wait(&cond_var, &mutex); ← release lock when wait
                                                acquire lock when being signaled
pthread_mutex_unlock(&mutex);
```

- Thread signaling another thread waiting on the condition variable:

```
pthread_mutex_lock(&mutex);
a = b;
pthread_cond_signal(&cond_var);
pthread_mutex_unlock(&mutex);
```

Takeaway

- Bounded-buffer problem
- Readers-writers problem
- Dining-philosophers problem
- Linux provides:
 - atomic integers
 - spinlocks
 - semaphores
 - reader-writer locks